

⌵ Importing Libraries

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from scipy import stats
from scipy.stats import probplot, shapiro

from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression, Ridge, Lasso
from sklearn.preprocessing import LabelEncoder, StandardScaler, MinMaxScaler
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score

import statsmodels.api as sm
from statsmodels.stats.outliers_influence import variance_inflation_factor
```

+ Code

+ Text

⌵ Loading Dataset

```
!wget "https://d2beiqkhq929f0.cloudfront.net/public_assets/assets/000/001/839/original/Jamboree_Admission.csv"

--2024-02-14 11:26:10-- https://d2beiqkhq929f0.cloudfront.net/public_assets/assets/000/001/839/original/Jamboree_Admission.csv
Resolving d2beiqkhq929f0.cloudfront.net (d2beiqkhq929f0.cloudfront.net)... 99.84.178.132, 99.84.178.93, 99.84.178.172, ...
Connecting to d2beiqkhq929f0.cloudfront.net (d2beiqkhq929f0.cloudfront.net)|99.84.178.132|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 16176 (16K) [text/plain]
Saving to: 'Jamboree_Admission.csv'



Jamboree_Admission. 100%[=====] 15.80K --.-KB/s in 0s

2024-02-14 11:26:11 (144 MB/s) - 'Jamboree_Admission.csv' saved [16176/16176]
```

⌵ Analysis of Dataset

```
df = pd.read_csv('Jamboree_Admission.csv')
```

```
df.head()
```

	Serial No.	GRE Score	TOEFL Score	University Rating	SOP	LOR	CGPA	Research	Chance of Admit	
0	1	337	118	4	4.5	4.5	9.65	1	0.92	
1	2	324	107	4	4.0	4.5	8.87	1	0.76	
2	3	316	104	3	3.0	3.5	8.00	1	0.72	
3	4	322	110	3	3.5	2.5	8.67	1	0.80	
4	5	314	103	2	2.0	3.0	8.21	0	0.65	

Next steps:

Generate code with df

 View recommended plots

```
df.shape
```

(500, 9)

```
df.columns
```

```
Index(['Serial No.', 'GRE Score', 'TOEFL Score', 'University Rating', 'SOP',
      'LOR ', 'CGPA', 'Research', 'Chance of Admit '],
      dtype='object')
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 500 entries, 0 to 499
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   Serial No.            500 non-null   int64
1   GRE Score             500 non-null   int64
2   TOEFL Score           500 non-null   int64
3   University Rating     500 non-null   int64
4   SOP                   500 non-null   float64
5   LOR                   500 non-null   float64
6   CGPA                  500 non-null   float64
7   Research              500 non-null   int64
8   Chance of Admit       500 non-null   float64
dtypes: float64(4), int64(5)
memory usage: 35.3 KB
```

## Basic Data Cleaning and Exploration

### Rename Columns

```
df.rename(columns = {'LOR ' : 'LOR', 'Chance of Admit ' : 'Chance of Admit'}, inplace = True)
```

### Column-wise Unique Entries

```
#find out unique entries
for i in df.columns:
    print(f"Unique entries for column {i:18} = {df[i].nunique()}")
```

```
Unique entries for column Serial No.      = 500
Unique entries for column GRE Score       = 49
Unique entries for column TOEFL Score     = 29
Unique entries for column University Rating = 5
Unique entries for column SOP             = 9
Unique entries for column LOR             = 9
Unique entries for column CGPA            = 184
Unique entries for column Research        = 2
Unique entries for column Chance of Admit = 61
```

### Conversion of Categorical Attributes to Category

```
categorical_columns = ['University Rating', 'SOP', 'LOR', 'Research']
for data in categorical_columns:
    df[data] = df[data].astype('category')
```

### Updating float64 to float32 columns

```
floating_columns = ['CGPA', 'Chance of Admit']
for i in floating_columns:
    df[i] = df[i].astype('float32')
```

### Removing the Unique Row Identifier

```
df = df.drop(columns='Serial No.')
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 500 entries, 0 to 499
Data columns (total 8 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   GRE Score             500 non-null   int64
1   TOEFL Score           500 non-null   int64
2   University Rating     500 non-null   category
3   SOP                   500 non-null   category
4   LOR                   500 non-null   category
5   CGPA                  500 non-null   float32
6   Research              500 non-null   category
```

```
7  Chance of Admit    500 non-null    float32
dtypes: category(4), float32(2), int64(2)
memory usage: 14.9 KB
```

Missing Values in Dataset

```
df.isnull().sum()

GRE Score      0
TOEFL Score    0
University Rating  0
SOP            0
LOR            0
CGPA           0
Research       0
Chance of Admit  0
dtype: int64
```

Duplicate Rows

```
df.duplicated().sum()

0
```

```
df.head()
```

	GRE Score	TOEFL Score	University Rating	SOP	LOR	CGPA	Research	Chance of Admit
0	337	118	4	4.5	4.5	9.65	1	0.92
1	324	107	4	4.0	4.5	8.87	1	0.76
2	316	104	3	3.0	3.5	8.00	1	0.72
3	322	110	3	3.5	2.5	8.67	1	0.80
4	314	103	2	2.0	3.0	8.21	0	0.65

Next steps:

Generate code with df

☒ View recommended plots

Statistical Summary

```
df.describe().T
```

	count	mean	std	min	25%	50%	75%	max
GRE Score	500.0	316.47200	11.295148	290.00	308.0000	317.00	325.00	340.00
TOEFL Score	500.0	107.19200	6.081868	92.00	103.0000	107.00	112.00	120.00
CGPA	500.0	8.57644	0.604813	6.80	8.1275	8.56	9.04	9.92
Chance of Admit	500.0	0.72174	0.141140	0.34	0.6300	0.72	0.82	0.97

```
df.describe(include='category').T
```

	count	unique	top	freq
University Rating	500.0	5.0	3.0	162.0
SOP	500.0	9.0	4.0	89.0
LOR	500.0	9.0	3.0	99.0
Research	500.0	2.0	1.0	280.0

Univariate Analysis

Distribution Plot for Categorical Variables

```
plt.figure(figsize = (12, 10))
```

```
plt.subplot(2,2,1)
sns.countplot(data=df, x='University Rating', hue = 'University Rating', legend= False, palette = 'dark:green')
plt.title('Distribution as per University rating', {'font':'serif', 'size':15,'weight':'bold'})

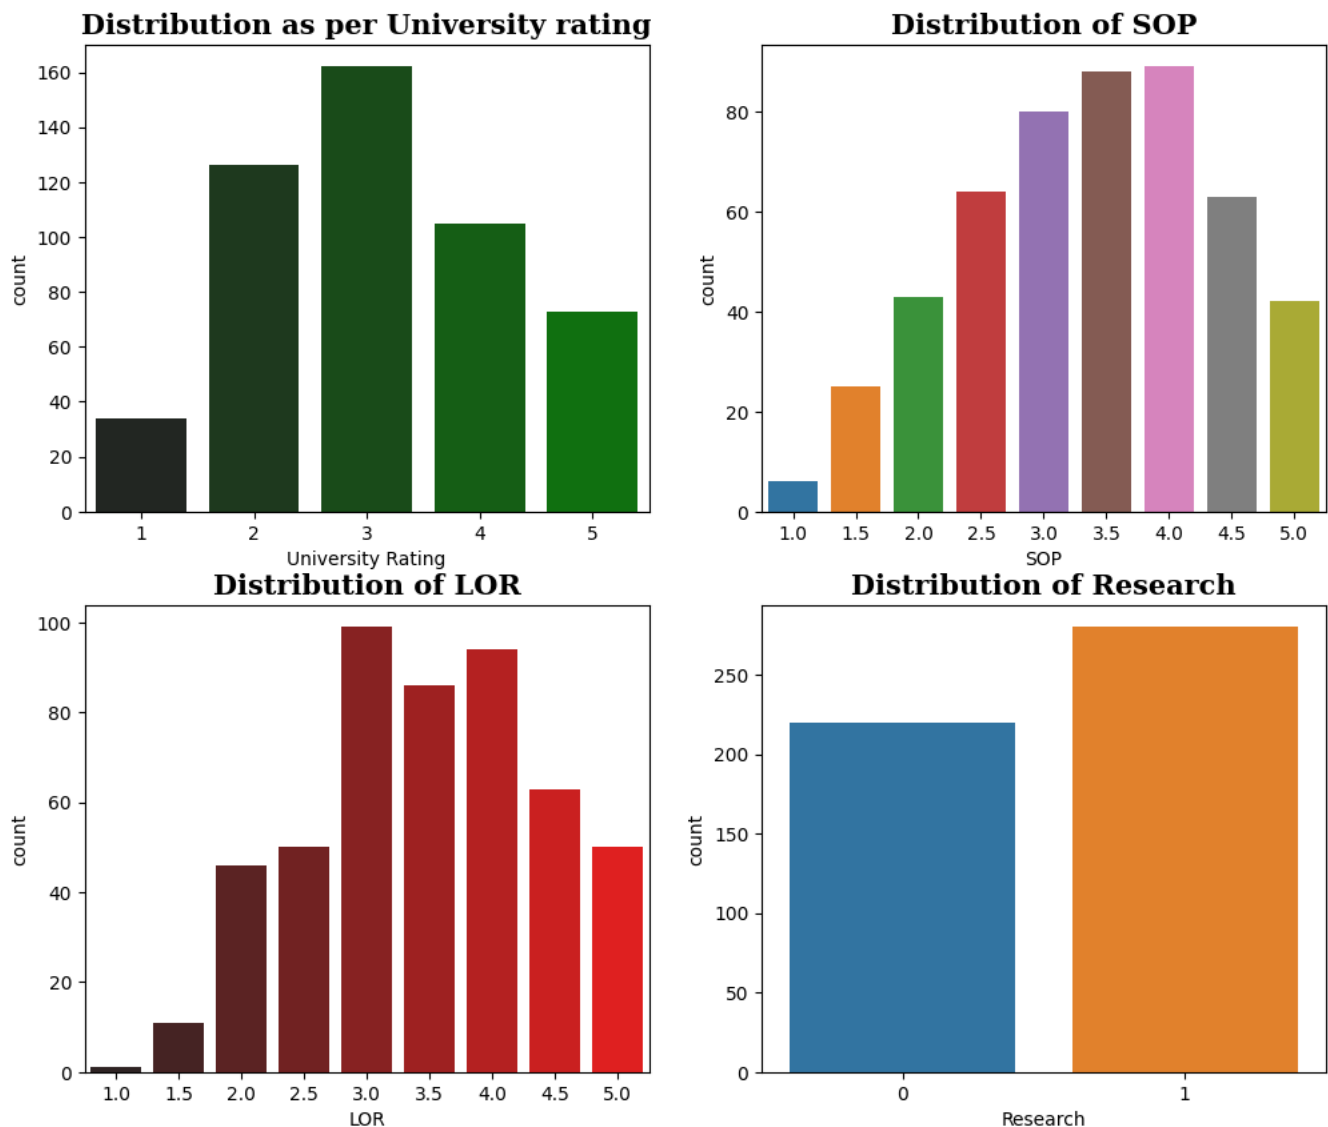
plt.subplot(2,2,2)
sns.countplot(data=df, x='SOP', hue = 'SOP', legend= False)
plt.title('Distribution of SOP', {'font':'serif', 'size':15,'weight':'bold'})

plt.subplot(2,2,3)
sns.countplot(data=df, x='LOR', hue = 'LOR', legend= False, palette = 'dark:red')
plt.title('Distribution of LOR', {'font':'serif', 'size':15,'weight':'bold'})

plt.subplot(2,2,4)
sns.countplot(data=df, x='Research', hue = 'Research', legend= False)
plt.title('Distribution of Research', {'font':'serif', 'size':15,'weight':'bold'})

plt.suptitle("Distribution of Categorical Variables", font='serif', size=20,weight='bold')
plt.show()
```

## Distribution of Categorical Variables



### Distribution Plot for Continuous Variables

```
plt.figure(figsize=(12,10))

plt.subplot(2,2,1)
sns.histplot(df['GRE Score'], kde = True, color = 'green', label = 'GRE Score')
plt.title("Distribution Plot for GRE Score", {'font':'serif', 'size':15, 'weight':'bold'})
```

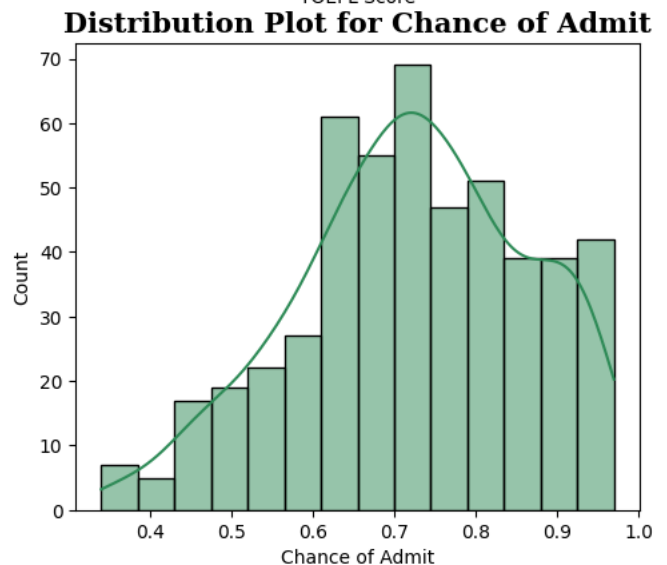



```
plt.subplot(2,2,2)
sns.histplot(df['TOEFL Score'], kde = True, color = 'red', label = 'TOEFL Score')
plt.title("Distribution Plot for TOEFL Score", {'font':'serif', 'size':15, 'weight':'bold'})

plt.subplot(2,2,3)
sns.histplot(df['CGPA'], kde = True, color = 'blue', label = 'CGPA')
plt.title("Distribution Plot for CGPA", {'font':'serif', 'size':15, 'weight':'bold'})

plt.subplot(2,2,4)
sns.histplot(df['Chance of Admit'], kde = True, color = 'seagreen', label = 'Chance of Admit')
plt.title("Distribution Plot for Chance of Admit", {'font':'serif', 'size':15, 'weight':'bold'})

plt.suptitle("Distribution Plots for Contineous Variables", font = 'serif', size = 20, weight = 'bold')
```

## Distribution Plots for Contineous Variables



### Bi-Variate Analysis

#### Relationship Between GRE Score, TOEFL Score and CGPA with Chance of Admit

```
plt.figure(figsize=(16,4))


plt.subplot(1,3,1)
sns.scatterplot(x=df['GRE Score'], y=df['Chance of Admit'])
```

```
plt.title("GRE Score and Chance of Admit", {'font':'serif', 'size':12, 'weight':'bold'})

plt.subplot(1,3,2)
sns.scatterplot(x=df['TOEFL Score'], y=df['Chance of Admit'])
plt.title("TOEFL Score and Chance of Admit", {'font':'serif', 'size':12, 'weight':'bold'})

plt.subplot(1,3,3)
sns.scatterplot(x=df['CGPA'], y=df['Chance of Admit'])
plt.title("CGPA and Chance of Admit", {'font':'serif', 'size':12, 'weight':'bold'})

plt.show()
```



### Relationship between Categorical columns and Chance of Admit

```
plt.figure(figsize=(14,10))

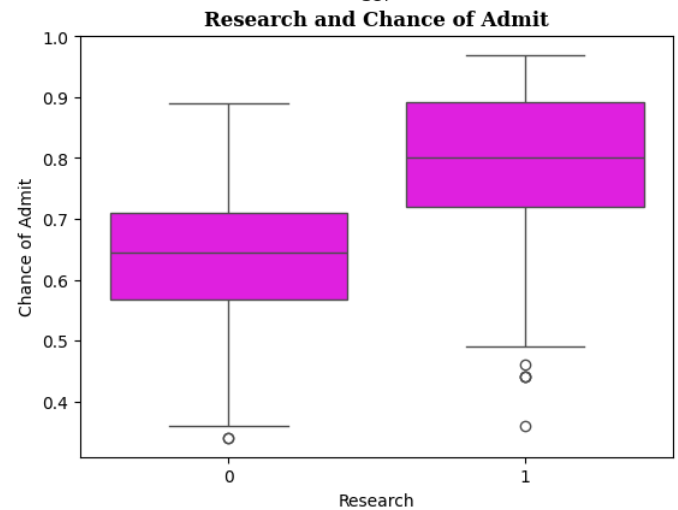
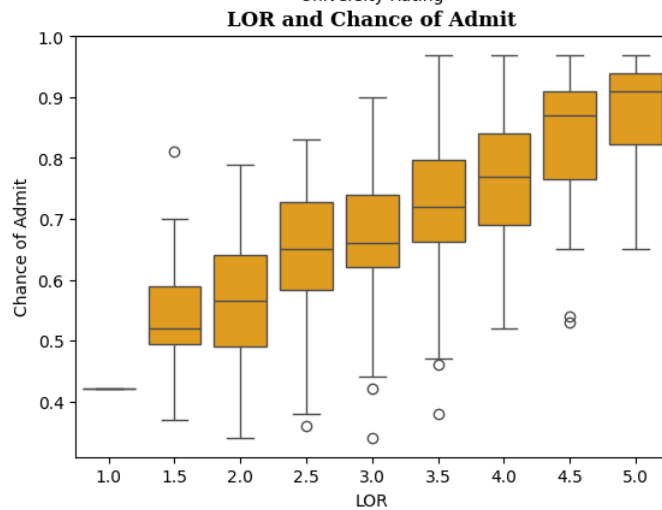
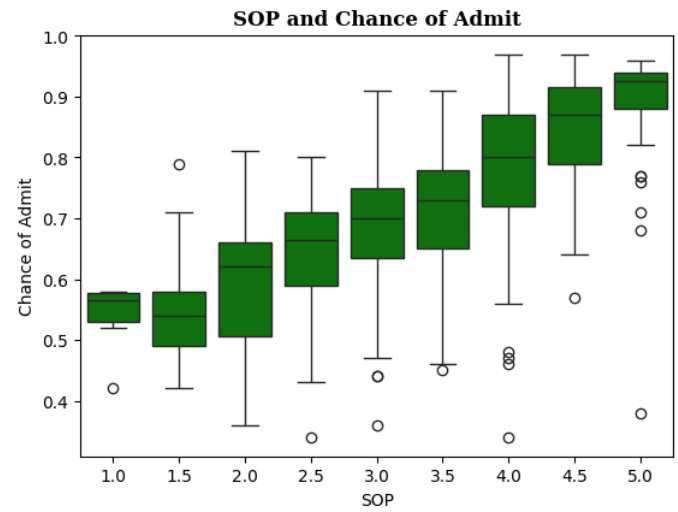

plt.subplot(2,2,1)
sns.boxplot(x= df['University Rating'],y= df['Chance of Admit'])
plt.title("University Rating and Chance of Admit", {'font':'serif', 'size':12, 'weight':'bold'})

plt.subplot(2,2,2)
sns.boxplot(x= df['SOP'],y= df['Chance of Admit'], color='green')
plt.title("SOP and Chance of Admit", {'font':'serif', 'size':12, 'weight':'bold'})

plt.subplot(2,2,3)
sns.boxplot(x= df['LOR'],y= df['Chance of Admit'], color='orange')
plt.title("LOR and Chance of Admit", {'font':'serif', 'size':12, 'weight':'bold'})

plt.subplot(2,2,4)
sns.boxplot(x= df['Research'],y= df['Chance of Admit'], color='magenta')
plt.title("Research and Chance of Admit", {'font':'serif', 'size':12, 'weight':'bold'})

plt.show()
```



### Multi-variate Analysis


```
plt.figure(figsize = (10, 6))

sns.heatmap(data = df.corr(), annot = True)

plt.title('Correlation between diffirent variables using heat map', font='serif', size=15, weight='bold')
#plt.xticks(rotation=45)

plt.show()
```

```
<ipython-input-25-26bf52663cf8>:3: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future version
sns.heatmap(data = df.corr(), annot = True)
```



## Outlier Detection

### Outlier detection by IQR method:

```
#Detecting IQR
def outlier_by_IQR(i):
    Q1 = np.quantile(df[i], 0.25)
    Q3 = np.quantile(df[i], 0.75)
    IQR = Q3 - Q1
    Lower_outlier = Q1 - 1.5*IQR
    Higher_outlier = Q3 + 1.5*IQR
    outliers = df.loc[(df[i] < Lower_outlier) | (df[i] > Higher_outlier)]
    print('Column :', i)
    print(f'Q1 : {Q1}')
    print(f'Q3 : {Q3}')
    print(f'IQR : {IQR}')
    print(f'Lower outlier : {Lower_outlier}')
    print(f'Upper outlier : {Higher_outlier}')
    print(f'Number of outliers : {outliers.shape[0]}')
    print('-----')

#detect outliers by IQR for contineous variables
numerical_columns = ['GRE Score', 'TOEFL Score', 'CGPA', 'Chance of Admit']
for i in numerical_columns:
    outlier_by_IQR(i)
```

```
Column : GRE Score
Q1 : 308.0
Q3 : 325.0
IQR : 17.0
Lower outlier : 282.5
Upper outlier : 350.5
Number of outliers : 0
-----
Column : TOEFL Score
Q1 : 103.0
Q3 : 112.0
IQR : 9.0
Lower outlier : 89.5
Upper outlier : 125.5
Number of outliers : 0
-----
Column : CGPA
```



```

Q1 : 8.127500057220459
Q3 : 9.039999961853027
IQR : 0.9124999046325684
Lower outlier : 6.7587502002716064
Upper outlier : 10.40874981880188
Number of outliers : 0
-----
Column : Chance of Admit
Q1 : 0.6299999952316284
Q3 : 0.8199999928474426
IQR : 0.1899999976158142
Lower outlier : 0.3449999988079071
Upper outlier : 1.104999989271164
Number of outliers : 2
-----

```

## ✓ Model Building

### *Data preparation for Model Building*

```

x = df.drop('Chance of Admit', axis=1)
y = df['Chance of Admit']

```

```

print("Shape of x: ", x.shape)
print("Shape of y: ", y.shape)

```

```

Shape of x: (500, 7)
Shape of y: (500,)

```

```

#split the data in test and train data
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size = 0.2, random_state = 42)

```

```

print("Shape of x_train: ", x_train.shape)
print("Shape of y_train: ", y_train.shape)

```

```

Shape of x_train: (400, 7)
Shape of y_train: (400,)

```

```

print("Shape of x_test: ", x_test.shape)
print("Shape of y_test: ", y_test.shape)

```

```

Shape of x_test: (100, 7)
Shape of y_test: (100,)

```

### *Transforming Categorical Columns by Label Encoding*

```

categorical_columns

```

```

['University Rating', 'SOP', 'LOR', 'Research']

```

```

#Transforming categorical columns in the train data and test data
Label_encoder = LabelEncoder()

```

```

#encode label in column-wise for train data
for i in categorical_columns:
    x_train[i] = Label_encoder.fit_transform(x_train[i])

```

```

#encode label in column-wise for test data
for i in categorical_columns:
    x_test[i] = Label_encoder.fit_transform(x_test[i])

```

### *Normalizing data by MinMaxScaler*

```

scaler = MinMaxScaler()

```

```

#Normalizing train data
x_train = pd.DataFrame(scaler.fit_transform(x_train), columns = x_train.columns)

```

x\_train.head()

	GRE Score	TOEFL Score	University Rating	SOP	LOR	CGPA	Research	
0	0.62	0.678571	0.50	0.625	0.714286	0.650641	1.0	
1	0.52	0.678571	0.75	0.750	1.000000	0.557692	0.0	
2	0.26	0.357143	0.50	0.625	0.428571	0.544872	0.0	
3	0.48	0.535714	0.25	0.375	0.714286	0.471154	0.0	
4	0.36	0.500000	0.50	0.625	0.285714	0.451923	1.0	

Next steps:

[Generate code with x\\_train](#)

☒ [View recommended plots](#)

x\_test.head()

	GRE Score	TOEFL Score	University Rating	SOP	LOR	CGPA	Research	
0	0.88	0.851852	0.75	0.750	0.625	0.894309	1.0	
1	0.48	0.555556	0.75	0.875	0.750	0.691057	1.0	
2	0.50	0.444444	0.25	0.250	0.375	0.126016	0.0	
3	0.44	0.592593	0.50	0.500	0.500	0.548780	0.0	
4	0.72	0.703704	0.50	0.625	0.500	0.695122	1.0	

Next steps:

[Generate code with x\\_test](#)

☒ [View recommended plots](#)

Linear Regression

Model Evaluation

```
model_LR = LinearRegression()
```

```
#fit the model in training data
model_LR.fit(x_train, y_train)
```

LinearRegression

LinearRegression()

Model Coefficients with Column Names

```
#bias
W0 = model_LR.intercept_
print(f"Intercept: {W0}")
```

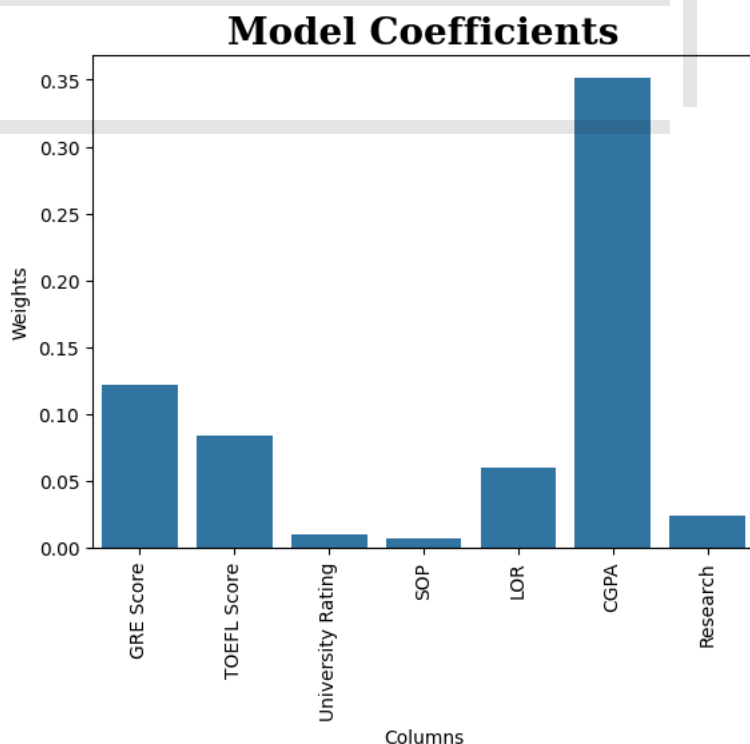
Intercept: 0.35558406163802325

```
#model coefficients
weights = pd.DataFrame({"Column" : x.columns, "Weight" : model_LR.coef_})
print(weights)
```

	Column	Weight
0	GRE Score	0.121722
1	TOEFL Score	0.083884
2	University Rating	0.010275
3	SOP	0.007255
4	LOR	0.060333
5	CGPA	0.351085
6	Research	0.024027

```
#Visualization of model coefficients
sns.barplot(x = x_train.columns, y = model_LR.coef_)
plt.title("Model Coefficients", font = 'serif', size= 20, weight= 'bold')
plt.xlabel("Columns")
plt.ylabel("Weights")
```

```
plt.xticks(rotation=90)
plt.show()
```



### Model Performance Evaluation

```
#reshaping of target train value
y_train = y_train.values.reshape(-1, 1)
```

```
#reshaping of target test value
y_test = y_test.values.reshape(-1, 1)
```

```
y_train.shape, y_test.shape
```

```
((400, 1), (100, 1))
```

```
#prediction value for train data
y_train_LR = model_LR.predict(x_train)
```

```
#prediction value for test data
y_test_LR = model_LR.predict(x_test)
```

```
def model_evaluation(y_actual, y_prediction, model, n, d):
```

```
    MAE = mean_absolute_error(y_actual, y_prediction)
```

```
    MSE = mean_squared_error(y_actual, y_prediction)
```

```
    RMSE = np.sqrt(MSE)
```

```
    R2 = r2_score(y_actual, y_prediction)
```

```
    Adjusted_R2 = 1 - ((1-R2)*(n-1)/(n - d - 1))
```

```
    return print(f"MAE: {round(MAE,2)}\nMSE: {round(MSE,3)}\nRMSE: {round(RMSE,2)}\nR2 score: {round(R2,2)}\nAdjusted R2: {round(Adjusted_R2,2)}")
```

```
#check the training data
```

```
print('Linear Regression Training Model\n')
```

```
model_evaluation(y_train, y_train_LR, model_LR, n = x_train.shape[0], d = x_train.shape[1])
```

```
#check the test data
```

```
print('\nLinear Regression Test Model\n')
```

```
model_evaluation(y_test, y_test_LR, model_LR, n = x_test.shape[0], d = x_test.shape[1])
```

```
Linear Regression Training Model
```

```
MAE: 0.04
```

```
MSE: 0.004
```

```
RMSE: 0.06
```

```
R2 score: 0.82
```

```
Adjusted R2: 0.82
```

Linear Regression Test Model

MAE: 0.05  
MSE: 0.004  
RMSE: 0.07  
R2 score: 0.79  
Adjusted R2: 0.77

### Ridge and Lasso regression

```
model_R = Ridge()
model_L = Lasso()
```

```
#fitting the model to training data
model_R.fit(x_train, y_train)
model_L.fit(x_train, y_train)
```

▼ Lasso  
Lasso()

```
#prediction for training and test data
y_train_R = model_R.predict(x_train)
y_test_R = model_R.predict(x_test)
```

```
y_train_L = model_L.predict(x_train)
y_test_L = model_L.predict(x_test)
```

```
# Evaluating Model Performance
print('Ridge Regression Training Model\n')
model_evaluation(y_train, y_train_R, model_R, n = x_train.shape[0], d=x_train.shape[1])
print('\n\nRidge Regression Test Model\n')
model_evaluation(y_test, y_test_R, model_R, n = x_test.shape[0], d=x_test.shape[1])
print('\n\nLasso Regression Training Model\n')
model_evaluation(y_train, y_train_L, model_L, n = x_train.shape[0], d=x_train.shape[1])
print('\n\nLasso Regression Test Model\n')
model_evaluation(y_test, y_test_L, model_L, n = x_test.shape[0], d=x_test.shape[1])
```

Ridge Regression Training Model

MAE: 0.04  
MSE: 0.004  
RMSE: 0.06  
R2 score: 0.82  
Adjusted R2: 0.82

Ridge Regression Test Model

MAE: 0.05  
MSE: 0.004  
RMSE: 0.06  
R2 score: 0.8  
Adjusted R2: 0.79

Lasso Regression Training Model

MAE: 0.11  
MSE: 0.02  
RMSE: 0.14  
R2 score: -0.0  
Adjusted R2: -0.02

Lasso Regression Test Model

MAE: 0.12  
MSE: 0.021  
RMSE: 0.14  
R2 score: -0.01  
Adjusted R2: -0.08

### Linear Regression from (Statsmodel library)

```
#add the constant term
x_sm = sm.add_constant(x_train)

#performing the ordinary least squares regression and fitting the model
results = sm.OLS(y_train, x_sm).fit()

# statistical summary of the model
print(results.summary())
```

OLS Regression Results						
=====						
Dep. Variable:	y	R-squared:	0.821			
Model:	OLS	Adj. R-squared:	0.818			
Method:	Least Squares	F-statistic:	257.0			
Date:	Wed, 14 Feb 2024	Prob (F-statistic):	3.41e-142			
Time:	11:26:16	Log-Likelihood:	561.91			
No. Observations:	400	AIC:	-1108.			
Df Residuals:	392	BIC:	-1076.			
Df Model:	7					
Covariance Type:	nonrobust					
=====						
	coef	std err	t	P> t	[0.025	0.975]
-----						
const	0.3556	0.010	36.366	0.000	0.336	0.375
GRE Score	0.1217	0.029	4.196	0.000	0.065	0.179
TOEFL Score	0.0839	0.026	3.174	0.002	0.032	0.136
University Rating	0.0103	0.017	0.611	0.541	-0.023	0.043
SOP	0.0073	0.020	0.357	0.721	-0.033	0.047
LOR	0.0603	0.016	3.761	0.000	0.029	0.092
CGPA	0.3511	0.034	10.444	0.000	0.285	0.417
Research	0.0240	0.007	3.231	0.001	0.009	0.039
=====						
Omnibus:	86.232	Durbin-Watson:	2.050			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	190.099			
Skew:	-1.107	Prob(JB):	5.25e-42			
Kurtosis:	5.551	Cond. No.	23.4			
=====						

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Testing the Assumptions of the Linear Regression model

1. Multicollinearity check by VIF score



```
def check_VIF(X_t):
    vif = pd.DataFrame()

    vif['Features'] = X_t.columns
    vif['VIF'] = [variance_inflation_factor(X_t.values, i) for i in range(X_t.shape[1])]
    vif['VIF'] = round(vif['VIF'], 2)
    vif = vif.sort_values(by = "VIF", ascending = False)
    return vif
```



```
X_t = pd.DataFrame(x_train, columns=x_train.columns)
check_VIF(X_t)
```

	Features	VIF
5	CGPA	39.76
0	GRE Score	31.20
1	TOEFL Score	26.76
3	SOP	18.57
4	LOR	11.01
2	University Rating	10.95
6	Research	3.36



```
#drop CGPA and again check VIF
X_t.drop(columns = ['CGPA'], inplace = True)
check_VIF(X_t)
```

	Features	VIF	
0	GRE Score	24.83	
1	TOEFL Score	24.22	
3	SOP	17.26	
2	University Rating	10.90	
4	LOR	10.15	
5	Research	3.36	



```
#drop GRE Score and again check VIF
X_t.drop(columns = ['GRE Score'], inplace = True)
check_VIF(X_t)
```

	Features	VIF	
2	SOP	17.07	
0	TOEFL Score	12.73	
1	University Rating	10.79	
3	LOR	10.09	
4	Research	2.99	



```
#drop SOP and again check VIF
X_t.drop(columns = ['SOP'], inplace = True)
check_VIF(X_t)
```

	Features	VIF	
0	TOEFL Score	10.51	
1	University Rating	9.33	
2	LOR	8.17	
3	Research	2.98	

```
#drop TOEFL Score and again check VIF
X_t.drop(columns = ['TOEFL Score'], inplace = True)
check_VIF(X_t)
```

	Features	VIF	
0	University Rating	7.19	
1	LOR	6.49	
2	Research	2.77	

```
#drop University Rating and again check VIF
X_t.drop(columns = ['University Rating'], inplace = True)
check_VIF(X_t)
```

	Features	VIF	
0	LOR	2.44	
1	Research	2.44	

2. Mean of Residuals


```
residuals = y_train.reshape(-1) - y_train_LR.reshape(-1)
mean_of_residuals = np.mean(residuals)

mean_of_residuals

-3.1780134079895106e-17
```

### 3. Linearity of Variables

```
sns.scatterplot(x = y_train_LR.reshape(-1), y= residuals)
plt.title('Residual Plot', font = 'serif', size= 15, weight='bold')
plt.xlabel('Predicted Values')
plt.ylabel('Residuals')
plt.axhline(y=0, color='r', linestyle='--')
plt.show();
```



### 4. Test for Homoscedasticity

```
errors = y_train.reshape(-1) - y_train_LR.reshape(-1)

sns.scatterplot(x = y_train_LR.reshape(-1, ), y= errors.reshape(-1, ))
sns.lineplot([0,0], color='red')
plt.title("Homoscedasticity Test", font='serif', size=15, weight='bold')
plt.xlabel("Predicted")
plt.ylabel("Errors")
plt.show()
```


## Homoscedasticity Test

### 5. Normality of Residuals

Histogram

```
sns.histplot(residuals, element = 'step', kde = True, color = 'red')
plt.title("Histogram of Residuals", font = 'serif', size = 15, weight = 'bold')

plt.show()
```



Q-Q plot

```
#Normality check
probplot(residuals, plot = plt, dist = 'norm')
plt.title('Q-Q plot of Residuals', weight = 'bold')
plt.show()
```

